



DEEP LEARNING: CIFAR10

***CHABCHOUB aymen - LAMA Graih Jules - MALONGA-NKOUNKOU Gautier - OUMESSAOUD Rabah***

# Presentation of Cifar 10

- CIFAR stands for: Canadian Institutes For Advanced Research
- collection of images that are commonly used to train machine learning and computer vision algorithms.
- The dataset is composed of **10 classes**:
  - cars, trucks, airplanes, ships
  - deer, dogs, cats, birds, horses, frogs

# The dataset in theory

- The CIFAR-10 dataset consists of 60000
- 32x32 colour images
- 6000 images per class
- 50000 training images
- 10000 test images
- The classes are completely mutually exclusive

# The dataset in practice

- **5 pickled files**
  - data
  - labels
  - label\_names
- **data:**
  - 10000 x 3072 numpy array
  - 10000 is the number of images
  - $3072 = 32 * 32 * 3$ (channels)
- **Labels:**
  - list of 10000 numbers between 0-9.
- **label\_names:**
  - string list with label names (size 10)

# Dataset: conclusion



- Small dataset: fits in the ram!
- Efficient for quick learning and fine tuning
- Classes are not skewed

# Linear models



- 3072 pixels for the input
  - activation function:  
linear
- 10 classes for the input with Softmax activation function
- optimizer: Adam

- 3072 pixels for the input
  - activation function:  
TanH
- 10 classes for the input with Softmax activation function
- optimizer: Adam

# Linear model v1



# Linear model v2





# Linear models: conclusion



- The model's predictions won't improve by just tuning the hyperparameters
- Adding layers allows predicting more complex functions
- Different optimizer allow faster model converging

# Feed forward models



- 3072 pixels for the input
  - activation function: reLu
  - 2 hidden layers (1040)
- 10 classes for the input with Softmax activation function
- optimizer: Adadelta

This model is close to our Linear model V2



# Feed forward v1





# Feed forward: regularizer 0.01





# Feed forward: regularizer 0.05





# Feed forward: regularizer 0.1



# Feed forward: regularizer 0.05 + dropout 0.1



# Feed forward: regularizer 0.05 + dropout 0.2 + BatchNormalization









# Feed forward models

- 3072 pixels for the input
  - activation function: reLu
  - 5 hidden layers
- 10 classes for the input with Softmax activation function
- optimizer: Adam



# Feed forward: no regularization





# Feed forward: regularization 0.1





# Feed forward: regularization 0.05





# Feed forward: regularization 0.1 + dropout 0.1





# Feed forward: regularization 0.1 + dropout 0.2



# Feed forward: regularization 0.1 + dropout 0.2 + lr 0.001





# Feed forward: regularization 0.1 + dropout 0.2 + 1 layer (6 dense)







Feed forward: regularization 0.1 + dropout 0.2  
+ 1 layer + lr 0.001





# Convolutional neural network

- 32,32,3 input shape
  - activation function: reLu
  - kernel\_size: 3,3
  - filters: 32
- 2 hidden layers: 1040 neurons
- 10 classes for the input with Softmax activation function
- optimizer: Adam



# CNN





# CNN + regularization: 0.1



# CNN + regularization: 0.1 + filters: 64





# CNN(2conv) - regularization: 0.1 - filters: 32/64





# CNN(3conv) - regularization: 0.1

## - filters: 32/64/128





# CNN(3conv) - regularization: 0.1 - filters: 32/64/128 - dropout 0.2









CNN(4conv) - regularization: 0.1  
- filters: 32/32/64/64 - dropout 0.2  
LR: 0.001





CNN(4conv) - regularization: 0.0001  
- filters: 32/32/64/64 - dropout 0.2  
LR: 0.001





CNN(4conv) - regularization: 0.001  
- filters: 32/32/64/64 - dropout 0.4  
LR: 0.001



CNN(4conv) - regularization: 0.001  
- filters: 32/32/64/64/128/128 - dropout 0.4  
LR: 0.001





CNN(4conv) - regularization: 0.001  
- filters: 32/32/64/64/128/128/256/256 -  
dropout 0.4  
LR: 0.001



# Residual neural network



- Convolution block
- Residual block
- 1 or 2 hidden layers
- 10 classes for the input with Softmax activation function
- optimizer: Adam

# Resnet: 10 residual layers



# Resnet: 15 residual layers



## L2 regu: 0.001







# Resnet: 15 residual layers

L2 regu: 0.001 + dropout 0.4 + 2 dense



# Resnet: 4 conv blocks + 12 residual layers

L2 regu: 0.001 + dropout 0.4 + 1 dense 512



# Resnet: 4 conv blocks + 12 residual layers

L2 regu: 0.001 + dropout 0.4 + 1 dense 1024

